

# POMDPs: Myths, Legends, and Reality

John Biechele-Speziale, Kushagra Kapoor, Jae Heo  
May 24, 2023

# Outline

What are POMDPs?

POMDPs are MDPs with sensors instead of direct observation of the current state.

- ▶ The sensors are some probability distribution of some observation given a certain state.

POMDPs are MDPs with sensors instead of direct observation of the current state.


- ▶ The sensors are some probability distribution of some observation given a certain state.
- ▶ Unlike traditional MDPs that map from states to actions, these map from belief states(observations) to actions.

POMDPs are MDPs with sensors instead of direct observation of the current state.

- ▶ The sensors are some probability distribution of some observation given a certain state.
- ▶ Unlike traditional MDPs that map from states to actions, these map from belief states(observations) to actions.
- ▶ Exact optimal solutions yield the optimal action for each possible belief state that minimizes our cost.

# POMDPs are MDPs with sensors instead of direct observation of the current state.

- ▶ The sensors are some probability distribution of some observation given a certain state.
- ▶ Unlike traditional MDPs that map from states to actions, these map from belief states(observations) to actions.
- ▶ Exact optimal solutions yield the optimal action for each possible belief state that minimizes our cost.



POMDPs are defined by a tuple  $(S, A, T, R, \Omega, O, \gamma)$

- ▶  $S$  is the set of states



POMDPs are defined by a tuple  $(S, A, T, R, \Omega, O, \gamma)$

- ▶  $S$  is the set of states
- ▶  $A$  is the set of actions

POMDPs are defined by a tuple  $(S, A, T, R, \Omega, O, \gamma)$

- ▶  $S$  is the set of states
- ▶  $A$  is the set of actions
- ▶  $T$  is the state transition probability  $P(s'|s, a)$

POMDPs are defined by a tuple  $(S, A, T, R, \Omega, O, \gamma)$

- ▶  $S$  is the set of states
- ▶  $A$  is the set of actions
- ▶  $T$  is the state transition probability  $P(s'|s, a)$
- ▶  $R$  is the reward function  $R(s, a)$

POMDPs are defined by a tuple  $(S, A, T, R, \Omega, O, \gamma)$

- ▶  $S$  is the set of states
- ▶  $A$  is the set of actions
- ▶  $T$  is the state transition probability  $P(s'|s, a)$
- ▶  $R$  is the reward function  $R(s, a)$
- ▶  $\Omega$  is the set of observations

# POMDPs are defined by a tuple $(S, A, T, R, \Omega, O, \gamma)$

- ▶  $S$  is the set of states
- ▶  $A$  is the set of actions
- ▶  $T$  is the state transition probability  $P(s'|s, a)$
- ▶  $R$  is the reward function  $R(s, a)$
- ▶  $\Omega$  is the set of observations
- ▶  $O$  is the set of conditional observation probabilities  $P(o|s') \vee P(o|s', a)$

# POMDPs are defined by a tuple $(S, A, T, R, \Omega, O, \gamma)$

- ▶  $S$  is the set of states
- ▶  $A$  is the set of actions
- ▶  $T$  is the state transition probability  $P(s'|s, a)$
- ▶  $R$  is the reward function  $R(s, a)$
- ▶  $\Omega$  is the set of observations
- ▶  $O$  is the set of conditional observation probabilities  $P(o|s') \vee P(o|s', a)$
- ▶  $\gamma$  is the discount factor.

# POMDPs are defined by a tuple $(S, A, T, R, \Omega, O, \gamma)$

- ▶  $S$  is the set of states
- ▶  $A$  is the set of actions
- ▶  $T$  is the state transition probability  $P(s'|s, a)$
- ▶  $R$  is the reward function  $R(s, a)$
- ▶  $\Omega$  is the set of observations
- ▶  $O$  is the set of conditional observation probabilities  $P(o|s') \vee P(o|s', a)$
- ▶  $\gamma$  is the discount factor.
- ▶ As usual, we want an optimal policy  $(\pi)$  that maximizes expected future reward.

POMDPs rely on a unique belief update step as a part of the algorithm.

► a



# Notes on theory, undecidability, and intractability



# POMDP Variants and Applications

Various methods of solving/approximating POMDPs are currently available/used.

- ▶ On Policy Methods
  - ▶
- ▶ Off Policy Methods
  - ▶

## Implementation Recommendations

My favorite package: POMDPs.jl is my favorite  
because it's easy to use: Installation

```
using Pkg; Pkg.add("POMDPs"); Pkg.add("QMDP");  
using POMDPs, QuickPOMDPs, POMDPTools, QMDP
```

My favorite package: POMDPs.jl is my favorite because it's easy to use: Definitions

```
m = QuickPOMDP(
    states = ["left", "right"],
    actions = ["left", "right", "listen"],
    observations = ["left", "right"],
    initialstate = Uniform(["left", "right"]),
    discount = 0.95,
    transition = function (s, a)
        if a == "listen"
            return Deterministic(s) # tiger stays behind the same
        else # a door is opened
            return Uniform(["left", "right"]) # reset
        end
    end,
    observation = function (s, a, sp)
        if a == "listen"
            if sp == "left"
                return SparseCat(["left", "right"], [0.85, 0.15])
            else
```

# Alternative Packages

- ▶ Finite-state Controllers using Branch and Bound
- ▶ pomdp
- ▶ pyPOMDP
- ▶ zmdp

Temp 1



# Temp 2

# Temp 3

# Temp 4

# Temp 5

# References I